


1

Introduction

Introduction

Parallel Multicore Architectures


- Increasingly widespread
- Increasingly dense
- Increasingly diverse
 - Specialized cores
 - Heterogeneity



Heterogeneous Parallel Platforms

Heterogeneous Association

- General purpose processor
- Specialized accelerator




Heterogeneous Parallel Platforms

Heterogeneous Association

- General purpose processor
- Specialized accelerator

Generalization



- Combination of various units
 - Latency-optimized cores
 - Throughput-optimized cores
 - Energy-optimized cores
- Distributed cores
 - Standalone GPUs
 - Intel Xeon Phi (MIC)
 - Intel Single-Chip Cloud (SCC)
- Integrated cores
 - Intel Haswell
 - AMD Fusion
 - nVidia Tegra
 - ARM big.LITTLE



Programming Models

How to Program these architectures?

- Multicore programming
 - pthreads, OpenMP, TBB, ...
- Accelerator programming
 - Consensus on OpenCL?
 - (Often) Pure offloading model
- Hybrid models?
 - Take advantage of all resources
 - Complex interactions



Work Needed at Multiple Levels

- Applications
 - Programming paradigm
 - BLAS kernels, FFT, ...
- Compilers
 - Languages
 - Code generation/optimization
- **Runtime systems**
 - Resources management
 - **Heterogeneous Task scheduling**
- Architecture
 - Memory interconnect

Heterogeneous Task Scheduling

Scheduling on platform equipped with accelerators

- Adapting to heterogeneity
 - Decide about tasks to offload
 - Decide about tasks to keep on CPU

Heterogeneous Task Scheduling

Scheduling on platform equipped with accelerators

- Adapting to heterogeneity
 - Decide about tasks to offload
 - Decide about tasks to keep on CPU
- Communicate with discrete accelerator board(s)
 - Send computation requests
 - Send data to be processed
 - Fetch results back
 - Expensive

Heterogeneous Task Scheduling

Scheduling on platform equipped with accelerators

- Adapting to heterogeneity
 - Decide about tasks to offload
 - Decide about tasks to keep on CPU
- Communicate with discrete accelerator board(s)
 - Send computation requests
 - Send data to be processed
 - Fetch results back
 - Expensive
- Decide about worthiness

2

StarPU Programming/Execution Models

Task Parallelism

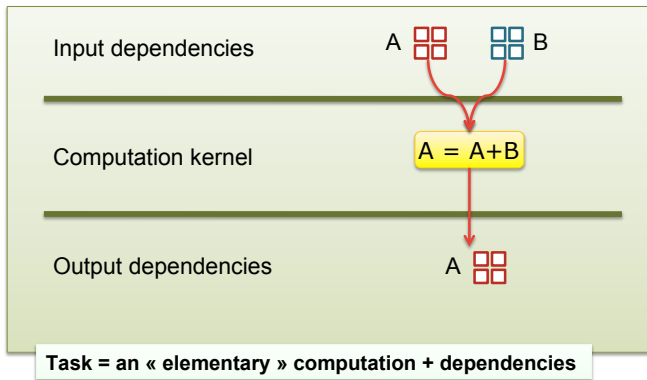
Principles

- Input dependencies
- Computation kernel
- Output dependencies

Task Parallelism

Principles

- Input dependencies
- Computation kernel
- Output dependencies



StarPU **Programming** Model: Sequential Task Flow

- Express parallelism. . .

StarPU **Programming** Model: Sequential Task Flow

- Express parallelism...
- ... using the natural program flow

StarPU **Programming** Model: Sequential Task Flow

- Express parallelism...
- ... using the natural program flow
- **Submit** tasks in the **sequential** flow of the program...

StarPU **Programming** Model: Sequential Task Flow



- Express parallelism. . .
- . . . using the natural program flow
- **Submit** tasks in the **sequential** flow of the program. . .
- . . . then let the runtime schedule the tasks **asynchronously**

Ex.: The Sequential **Task-Based** Cholesky Decomposition

```
for (j = 0; j < N; j++) {  
    POTRF (RW,A[j][j]);  
    for (i = j+1; i < N; i++)  
        TRSM (RW,A[i][j], R,A[j][j]);  
    for (i = j+1; i < N; i++) {  
        SYRK (RW,A[i][i], R,A[i][j]);  
        for (k = j+1; k < i; k++)  
            GEMM (RW,A[i][k],  
                R,A[i][j], R,A[k][j]);  
    }  
}  
__wait__();
```


Ex.: The Sequential **Task-Based** Cholesky Decomposition

```
for (j = 0; j < N; j++) {  
    POTRF (RW,A[j][j]);  
    for (i = j+1; i < N; i++)  
        TRSM (RW,A[i][j], R,A[j][j]);  
    for (i = j+1; i < N; i++) {  
        SYRK (RW,A[i][i], R,A[i][j]);  
        for (k = j+1; k < i; k++)  
            GEMM (RW,A[i][k],  
                  R,A[i][j], R,A[k][j]);  
    }  
}  
__wait__();
```




Ex.: The Sequential **Task-Based** Cholesky Decomposition

```
for (j = 0; j < N; j++) {  
    POTRF (RW,A[j][j]);  
    for (i = j+1; i < N; i++)  
        TRSM (RW,A[i][j], R,A[j][j]);  
    for (i = j+1; i < N; i++) {  
        SYRK (RW,A[i][i], R,A[i][j]);  
        for (k = j+1; k < i; k++)  
            GEMM (RW,A[i][k],  
                R,A[i][j], R,A[k][j]);  
    }  
}  
__wait__();
```





Ex.: The Sequential **Task-Based** Cholesky Decomposition

```
for (j = 0; j < N; j++) {  
    POTRF (RW,A[j][j]);  
    for (i = j+1; i < N; i++)  
        TRSM (RW,A[i][j], R,A[j][j]);  
    for (i = j+1; i < N; i++) {  
        SYRK (RW,A[i][i], R,A[i][j]);  
        for (k = j+1; k < i; k++)  
            GEMM (RW,A[i][k],  
                R,A[i][j], R,A[k][j]);  
    }  
}  
__wait__();
```






Ex.: The Sequential **Task-Based** Cholesky Decomposition

```
for (j = 0; j < N; j++) {  
    POTRF (RW,A[j][j]);  
    for (i = j+1; i < N; i++)  
        TRSM (RW,A[i][j], R,A[j][j]);  
    for (i = j+1; i < N; i++) {  
        SYRK (RW,A[i][i], R,A[i][j]);  
        for (k = j+1; k < i; k++)  
            GEMM (RW,A[i][k],  
                R,A[i][j], R,A[k][j]);  
    }  
}  
__wait__();
```






Ex.: The Sequential **Task-Based** Cholesky Decomposition

```
for (j = 0; j < N; j++) {  
    POTRF (RW,A[j][j]);  
    for (i = j+1; i < N; i++)  
        TRSM (RW,A[i][j], R,A[j][j]);  
    for (i = j+1; i < N; i++) {  
        SYRK (RW,A[i][i], R,A[i][j]);  
        for (k = j+1; k < i; k++)  
            GEMM (RW,A[i][k],  
                R,A[i][j], R,A[k][j]);  
    }  
}  
__wait__();
```






Ex.: The Sequential **Task-Based** Cholesky Decomposition

```
for (j = 0; j < N; j++) {  
    POTRF (RW,A[j][j]);  
    for (i = j+1; i < N; i++)  
        TRSM (RW,A[i][j], R,A[j][j]);  
    for (i = j+1; i < N; i++) {  
        SYRK (RW,A[i][i], R,A[i][j]);  
        for (k = j+1; k < i; k++)  
            GEMM (RW,A[i][k],  
                R,A[i][j], R,A[k][j]);  
    }  
}  
__wait__();
```





Ex.: The Sequential **Task-Based** Cholesky Decomposition

```
for (j = 0; j < N; j++) {  
    POTRF (RW,A[j][j]);  
    for (i = j+1; i < N; i++)  
        TRSM (RW,A[i][j], R,A[j][j]);  
    for (i = j+1; i < N; i++) {  
        SYRK (RW,A[i][i], R,A[i][j]);  
        for (k = j+1; k < i; k++)  
            GEMM (RW,A[i][k],  
                 R,A[i][j], R,A[k][j]);  
    }  
}  
__wait__();
```





Ex.: The Sequential **Task-Based** Cholesky Decomposition

```
for (j = 0; j < N; j++) {  
    POTRF (RW,A[j][j]);  
    for (i = j+1; i < N; i++)  
        TRSM (RW,A[i][j], R,A[j][j]);  
    for (i = j+1; i < N; i++) {  
        SYRK (RW,A[i][i], R,A[i][j]);  
        for (k = j+1; k < i; k++)  
            GEMM (RW,A[i][k],  
                R,A[i][j], R,A[k][j]);  
    }  
}  
__wait__();
```






Ex.: The Sequential **Task-Based** Cholesky Decomposition

```
for (j = 0; j < N; j++) {  
    POTRF (RW,A[j][j]);  
    for (i = j+1; i < N; i++)  
        TRSM (RW,A[i][j], R,A[j][j]);  
    for (i = j+1; i < N; i++) {  
        SYRK (RW,A[i][i], R,A[i][j]);  
        for (k = j+1; k < i; k++)  
            GEMM (RW,A[i][k],  
                R,A[i][j], R,A[k][j]);  
    }  
}  
__wait__();
```






Ex.: The Sequential **Task-Based** Cholesky Decomposition

```
for (j = 0; j < N; j++) {  
    POTRF (RW,A[j][j]);  
    for (i = j+1; i < N; i++)  
        TRSM (RW,A[i][j], R,A[j][j]);  
    for (i = j+1; i < N; i++) {  
        SYRK (RW,A[i][i], R,A[i][j]);  
        for (k = j+1; k < i; k++)  
            GEMM (RW,A[i][k],  
                R,A[i][j], R,A[k][j]);  
    }  
}  
__wait__();
```






Ex.: The Sequential **Task-Based** Cholesky Decomposition

```
for (j = 0; j < N; j++) {  
    POTRF (RW,A[j][j]);  
    for (i = j+1; i < N; i++)  
        TRSM (RW,A[i][j], R,A[j][j]);  
    for (i = j+1; i < N; i++) {  
        SYRK (RW,A[i][i], R,A[i][j]);  
        for (k = j+1; k < i; k++)  
            GEMM (RW,A[i][k],  
                R,A[i][j], R,A[k][j]);  
    }  
}  
__wait__();
```





Ex.: The Sequential **Task-Based** Cholesky Decomposition

```
for (j = 0; j < N; j++) {  
    POTRF (RW,A[j][j]);  
    for (i = j+1; i < N; i++)  
        TRSM (RW,A[i][j], R,A[j][j]);  
    for (i = j+1; i < N; i++) {  
        SYRK (RW,A[i][i], R,A[i][j]);  
        for (k = j+1; k < i; k++)  
            GEMM (RW,A[i][k],  
                R,A[i][j], R,A[k][j]);  
    }  
}  
__wait__();
```



Ex.: The Sequential **Task-Based** Cholesky Decomposition

```
for (j = 0; j < N; j++) {  
    POTRF (RW,A[j][j]);  
    for (i = j+1; i < N; i++)  
        TRSM (RW,A[i][j], R,A[j][j]);  
    for (i = j+1; i < N; i++) {  
        SYRK (RW,A[i][i], R,A[i][j]);  
        for (k = j+1; k < i; k++)  
            GEMM (RW,A[i][k],  
                R,A[i][j], R,A[k][j]);  
    }  
}  
__wait__();
```

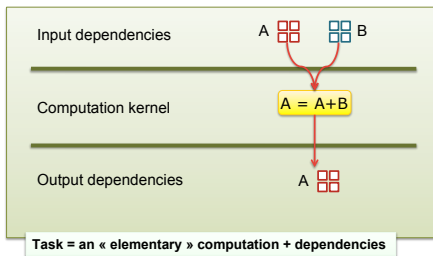


Task Relationships

Abstract Application Structure

Task Relationships



Abstract Application Structure



Task Relationships

Abstract Application Structure

- Directed Acyclic Graph (DAG)




StarPU **Execution** Model: Task Scheduling

Mapping the graph of tasks (DAG) on the hardware

StarPU **Execution** Model: Task Scheduling

Mapping the graph of tasks (DAG) on the hardware


- Allocating computing resources




StarPU **Execution** Model: Task Scheduling

Mapping the graph of tasks (DAG) on the hardware

- Allocating computing resources
- Enforcing dependency constraints
- Handling data transfers



A Single DAG for Multiple Schedules, Platforms



StarPU in a Nutshell

Rationale

- Implement the sequential task flow programming model
- Map computations on heterogeneous computing units

StarPU in a Nutshell

Rationale

- Implement the sequential task flow programming model
- Map computations on heterogeneous computing units

Programming Model

- Task
- Data
- Relationships
 - Task \leftrightarrow Task
 - Task \leftrightarrow Data

StarPU in a Nutshell

Rationale

- Implement the sequential task flow programming model
- Map computations on heterogeneous computing units

Programming Model




- Task
- Data
- Relationships
 - Task \leftrightarrow Task
 - Task \leftrightarrow Data

Runtime System




- Heterogeneous Task scheduling
- Application Programming Interface (Library)

What StarPU can do for You?




What StarPU does for You: **Heterogeneous Task Scheduling**



What StarPU does for You: **Heterogeneous Task Scheduling**





What StarPU does for You: **Heterogeneous Task Scheduling**






What StarPU does for You:

What StarPU does for You: **Data Transfers**

What StarPU does for You: **Data Transfers**



What StarPU does for You: **Data Transfers**



3

Programming with StarPU

Basic Example: Scaling a Vector

```
1 float factor = 3.14;  
2 float vector[NX];
```

Basic Example: Scaling a Vector

```
1 float factor = 3.14;  
2 float vector[NX];  
3 starpu_data_handle_t vector_handle;
```

Basic Example: Scaling a Vector

```
1 float factor = 3.14;
2 float vector[NX];
3 starpu_data_handle_t vector_handle;
4
5 /* ... fill vector ... */
6
7 starpu_vector_data_register(&vector_handle, 0,
8                             (uintptr_t)vector, NX, sizeof(vector[0]));
```

Basic Example: Scaling a Vector

```
1 float factor = 3.14;
2 float vector[NX];
3 starpu_data_handle_t vector_handle;
4
5 /* ... fill vector ... */
6
7 starpu_vector_data_register(&vector_handle, 0,
8                             (uintptr_t)vector, NX, sizeof(vector[0]));
9
10 starpu_task_insert(
11     &scal_cl,
12     STARPU_RW, vector_handle,
13     STARPU_VALUE, &factor, sizeof(factor),
14     0);
```

Basic Example: Scaling a Vector

```
1 float factor = 3.14;
2 float vector[NX];
3 starpu_data_handle_t vector_handle;
4
5 /* ... fill vector ... */
6
7 starpu_vector_data_register(&vector_handle, 0,
8                             (uintptr_t)vector, NX, sizeof(vector[0]));
9
10 starpu_task_insert(
11     &scal_cl,
12     STARPU_RW, vector_handle,
13     STARPU_VALUE, &factor, sizeof(factor),
14     0);
15
16 starpu_task_wait_for_all();
```

Basic Example: Scaling a Vector

```
1 float factor = 3.14;
2 float vector[NX];
3 starpu_data_handle_t vector_handle;
4
5 /* ... fill vector ... */
6
7 starpu_vector_data_register(&vector_handle, 0,
8                             (uintptr_t)vector, NX, sizeof(vector[0]));
9
10 starpu_task_insert(
11     &scal_cl,
12     STARPU_RW, vector_handle,
13     STARPU_VALUE, &factor, sizeof(factor),
14     0);
15
16 starpu_task_wait_for_all();
17 starpu_data_unregister(vector_handle);
18
19 /* ... display vector ... */
```

Terminology

- Codelet
- Task
- Data handle

Definition: A Codelet

A **Codelet**...

- ... relates an abstract computation kernel to its implementation(s)
- ... can be instantiated into one or more **tasks**
- ... defines characteristics common to a set of **tasks**

Definition: A Codelet

A Codelet...

- ... relates an abstract computation kernel to its implementation(s)
- ... can be instantiated into one or more **tasks**
- ... defines characteristics common to a set of **tasks**

Codelet
scal_cl



Definition: A Codelet

A Codelet...

- ... relates an abstract computation kernel to its implementation(s)
- ... can be instantiated into one or more **tasks**
- ... defines characteristics common to a set of **tasks**


Codelet
scal_cl



Definition: A Codelet

A Codelet...


- ... relates an abstract computation kernel to its implementation(s)
- ... can be instantiated into one or more **tasks**
- ... defines characteristics common to a set of **tasks**



Definition: A Codelet

A Codelet...


- ... relates an abstract computation kernel to its implementation(s)
- ... can be instantiated into one or more **tasks**
- ... defines characteristics common to a set of **tasks**



Definition: A Codelet

A Codelet...

- ... relates an abstract computation kernel to its implementation(s)
- ... can be instantiated into one or more **tasks**
- ... defines characteristics common to a set of **tasks**



Definition: A Task

A **Task**...

- ... is an instantiation of a **Codelet**
- ... atomically executes a kernel from its beginning to its end
- ... receives some input
- ... produces some output

Definition: A Task

A **Task**...

- ... is an instantiation of a **Codelet**
- ... atomically executes a kernel from its beginning to its end
- ... receives some input
- ... produces some output


Codelet
scal_cl



Definition: A Task

A Task...


- ... is an instantiation of a **Codelet**
- ... atomically executes a kernel from its beginning to its end
- ... receives some input
- ... produces some output



Definition: A Task

A **Task**...


- ... is an instantiation of a **Codelet**
- ... atomically executes a kernel from its beginning to its end
- ... receives some input
- ... produces some output



Definition: A Task

A **Task**...


- ... is an instantiation of a **Codelet**
- ... atomically executes a kernel from its beginning to its end
- ... receives some input
- ... produces some output



Definition: A Task

A Task...

- ... is an instantiation of a **Codelet**
- ... atomically executes a kernel from its beginning to its end
- ... receives some input
- ... produces some output



Definition: A Task

A **Task**...


- ... is an instantiation of a **Codelet**
- ... atomically executes a kernel from its beginning to its end
- ... receives some input
- ... produces some output



Definition: A Task

A Task...


- ... is an instantiation of a **Codelet**
- ... atomically executes a kernel from its beginning to its end
- ... receives some input
- ... produces some output



Definition: A Task

A Task...

- ... is an instantiation of a **Codelet**
- ... atomically executes a kernel from its beginning to its end
- ... receives some input
- ... produces some output




Task 1 is running

Definition: A Task

A Task...

- ... is an instantiation of a **Codelet**
- ... atomically executes a kernel from its beginning to its end
- ... receives some input
- ... produces some output



Definition: A Data Handle

A Data Handle...

- ... designates a piece of data managed by StarPU
- ... is typed (vector, matrix, etc.)
- ... can be passed as input/output for a **Task**

Elementary API

- Initializing/Ending a StarPU session
- Declaring a codelet
- Declaring and Managing Data
- Writing a Kernel Function
- Submitting a task
- Waiting for submitted tasks

Initializing a StarPU Session

- `starpu_init(struct starpu_conf *configuration)`

Initializing a StarPU Session

- `starpu_init(struct starpu_conf *configuration)`
 - The `struct starpu_conf` can be used to configure StarPU settings
 - Specify NULL for default settings

Initializing a StarPU Session

- **starpu_init**(**struct starpu_conf** *configuration)
 - The **struct starpu_conf** can be used to configure StarPU settings
 - Specify NULL for default settings

```
1 #include <starpu.h>
2
3 int ret = starpu_init(NULL);
4
5 if (ret == 0) {
6     printf("StarPU_successfully_initialized\n");
7 } else {
8     fprintf(stderr, "StarPU_initialization_failed\n");
9     exit(1);
10 }
11
12 /* StarPU is ready */
13 ...
```

Ending a StarPU Session

- `starpu_shutdown()`

Ending a StarPU Session

- `starpu_shutdown()`

```
1 ...  
2 starpu_shutdown() ;  
3  
4 /* StarPU is terminated */  
5 ...
```

Declaring a Codelet

Define a **struct** `starpu_codelet`

```
1 struct starpu_codelet scal_cl = {  
2     ...  
3 };
```


Declaring a Codelet

Define a **struct** `starpu_codelet`

- Plug the kernel function
 - Here: `scal_cpu_func`

```
1 struct starpu_codelet scal_cl = {  
2     .cpu_func = { scal_cpu_func, NULL },  
3     ...  
4 };
```

Declaring a Codelet

Define a **struct** `starpu_codelet`

- Plug the kernel function
 - Here: `scal_cpu_func`
- Declare the number of data pieces used by the kernel
 - Here: A single vector

```
1 struct starpu_codelet scal_cl = {  
2     .cpu_func = { scal_cpu_func, NULL },  
3     .nbuffers = 1,  
4     ...  
5 };
```

Declaring a Codelet

Define a **struct starpu_codelet**

- Plug the kernel function
 - Here: `scal_cpu_func`
- Declare the number of data pieces used by the kernel
 - Here: A single vector
- Declare how the kernel accesses the piece of data
 - Here: The vector is scaled in-place, thus R/W

```
1 struct starpu_codelet scal_cl = {  
2     .cpu_func = { scal_cpu_func, NULL },  
3     .nbuffers = 1,  
4     .modes = { STARPU_RW },  
5 };
```

Declaring and Managing Data

Put data under StarPU control

Declaring and Managing Data

Put data under StarPU control

- Initialize a piece of data

```
1 float vector[NX];  
2 /* ... fill data ... */
```

Declaring and Managing Data

Put data under StarPU control

- Initialize a piece of data
- Register the piece of data and get a handle
 - The vector is now under StarPU control

```
1 float vector[NX];  
2 /* ... fill data ... */  
3  
4 starpu_data_handle_t vector_handle;  
5 starpu_vector_data_register(&vector_handle, 0,  
6                             (uintptr_t)vector, NX, sizeof(vector[0]));
```

Declaring and Managing Data

Put data under StarPU control

- Initialize a piece of data
- Register the piece of data and get a handle
 - The vector is now under StarPU control
- Use data through the handle

```
1 float vector[NX];  
2 /* ... fill data ... */  
3  
4 starpu_data_handle_t vector_handle;  
5 starpu_vector_data_register(&vector_handle, 0,  
6                             (uintptr_t)vector, NX, sizeof(vector[0]));  
7  
8 /* ... use the vector through the handle ... */
```

Declaring and Managing Data

Put data under StarPU control

- Initialize a piece of data
- Register the piece of data and get a handle
 - The vector is now under StarPU control
- Use data through the handle
- Unregister the piece of data
 - The handle is destroyed
 - **The vector is now back under user control**

```
1 float vector[NX];  
2 /* ... fill data ... */  
3  
4 starpu_data_handle_t vector_handle;  
5 starpu_vector_data_register(&vector_handle, 0,  
6                             (uintptr_t)vector, NX, sizeof(vector[0]));  
7  
8 /* ... use the vector through the handle ... */  
9  
10 starpu_data_unregister(vector_handle);
```


Writing a Kernel Function

- Every kernel function has the same C prototype

```
1 void scal_cpu_func(void *buffers [], void *cl_arg) {  
2     ...  
3 }
```

Writing a Kernel Function

- Every kernel function has the same C prototype
- Retrieve the vector's handle

```
1 void scal_cpu_func(void *buffers[], void *cl_arg) {  
2     struct starpu_vector_interface *vector_handle = buffers[0];  
3  
4     ...  
5 }
```

Writing a Kernel Function

- Every kernel function has the same C prototype
- Retrieve the vector's handle
- Get vector's number of elements and base pointer

```
1 void scal_cpu_func(void *buffers[], void *cl_arg) {  
2     struct starpu_vector_interface *vector_handle = buffers[0];  
3  
4     unsigned n      = STARPU_VECTOR_GET_NX(vector_handle);  
5     float *vector = STARPU_VECTOR_GET_PTR(vector_handle);  
6  
7     ...  
8 }
```

Writing a Kernel Function

- Every kernel function has the same C prototype
- Retrieve the vector's handle
- Get vector's number of elements and base pointer
- Get the scaling factor as inline argument

```
1 void scal_cpu_func(void *buffers[], void *cl_arg) {  
2     struct starpu_vector_interface *vector_handle = buffers[0];  
3  
4     unsigned n      = STARPU_VECTOR_GET_NX(vector_handle);  
5     float *vector = STARPU_VECTOR_GET_PTR(vector_handle);  
6  
7     float *ptr_factor = cl_arg;  
8  
9     ...  
10 }
```

Writing a Kernel Function

- Every kernel function has the same C prototype
- Retrieve the vector's handle
- Get vector's number of elements and base pointer
- Get the scaling factor as inline argument
- Compute the vector scaling

```
1 void scal_cpu_func(void *buffers[], void *cl_arg) {  
2     struct starpu_vector_interface *vector_handle = buffers[0];  
3  
4     unsigned n      = STARPU_VECTOR_GET_NX(vector_handle);  
5     float *vector = STARPU_VECTOR_GET_PTR(vector_handle);  
6  
7     float *ptr_factor = cl_arg;  
8  
9     unsigned i;  
10    for (i = 0; i < n; i++)  
11        vector[i] *= *ptr_factor;  
12 }
```

Submitting a task

The `starpu_task_insert` call

- **Inserts** a task in the StarPU DAG

Submitting a task

The `starpu_task_insert` call

- **Inserts** a task in the StarPU DAG

Arguments

- The codelet structure

```
1 starpu_task_insert(&scal_cl  
2                  ...);
```

Submitting a task

The `starpu_task_insert` call

- **Inserts** a task in the StarPU DAG

Arguments

- The codelet structure
- The StarPU-managed data

```
1 starpu_task_insert(&scal_cl ,  
2                   STARPU_RW, vector_handle ,  
3                   ... ) ;
```


Submitting a task

The `starpu_task_insert` call

- **Inserts** a task in the StarPU DAG

Arguments

- The codelet structure
- The StarPU-managed data
- The small-size inline data

```
1 starpu_task_insert(&scal_cl ,  
2                     STARPU_RW, vector_handle ,  
3                     STARPU_VALUE, &factor , sizeof(factor) ,  
4                     ... ) ;
```

Submitting a task

The `starpu_task_insert` call

- **Inserts** a task in the StarPU DAG

Arguments

- The codelet structure
- The StarPU-managed data
- The small-size inline data
- 0 to mark the end of arguments

```
1 starpu_task_insert(&scal_cl ,  
2                     STARPU_RW, vector_handle ,  
3                     STARPU_VALUE, &factor , sizeof(factor) ,  
4                     0);
```

Submitting a task

The `starpu_task_insert` call

- **Inserts** a task in the StarPU DAG

Arguments

- The codelet structure
- The StarPU-managed data
- The small-size inline data
- 0 to mark the end of arguments

Notes

- The task is submitted non-blockingly

Submitting a task

The `starpu_task_insert` call

- **Inserts** a task in the StarPU DAG

Arguments

- The codelet structure
- The StarPU-managed data
- The small-size inline data
- 0 to mark the end of arguments

Notes

- The task is submitted non-blockingly
- Dependencies are enforced with previously submitted tasks' data. . .

Submitting a task

The `starpu_task_insert` call

- **Inserts** a task in the StarPU DAG

Arguments

- The codelet structure
- The StarPU-managed data
- The small-size inline data
- 0 to mark the end of arguments

Notes

- The task is submitted non-blockingly
- Dependencies are enforced with previously submitted tasks' data. . .
- . . . following the natural order of the program

Submitting a task

The `starpu_task_insert` call

- **Inserts** a task in the StarPU DAG

Arguments

- The codelet structure
- The StarPU-managed data
- The small-size inline data
- 0 to mark the end of arguments

Notes

- The task is submitted non-blockingly
- Dependencies are enforced with previously submitted tasks' data. . .
- . . . following the **natural** order of the program
- This is the **Sequential Task Flow Paradigm**

Waiting for Submitted Task Completion

- Tasks are submitted non-blockingly

Waiting for Submitted Task Completion

- Tasks are submitted non-blockingly

```
1  /* non-blocking task submits */  
2  starpu_task_insert (...);  
3  starpu_task_insert (...);  
4  starpu_task_insert (...);  
5  ...
```


Waiting for Submitted Task Completion

- Tasks are submitted non-blockingly
- Wait for all submitted tasks to complete their work

```
1  /* non-blocking task submits */  
2  starpu_task_insert (...);  
3  starpu_task_insert (...);  
4  starpu_task_insert (...);  
5  ...
```

Waiting for Submitted Task Completion

- Tasks are submitted non-blockingly
- Wait for all submitted tasks to complete their work

```
1  /* non-blocking task submits */
2  starpu_task_insert (...);
3  starpu_task_insert (...);
4  starpu_task_insert (...);
5  ...
6
7  /* wait for all task submitted so far */
8  starpu_task_wait_for_all();
```

Basic Example: Scaling a Vector

```
1 float factor = 3.14;
2 float vector[NX];
3 starpu_data_handle_t vector_handle;
4
5 /* ... fill vector ... */
6
7 starpu_vector_data_register(&vector_handle, 0,
8                             (uintptr_t)vector, NX, sizeof(vector[0]));
9
10 starpu_task_insert(
11     &scal_cl,
12     STARPU_RW, vector_handle,
13     STARPU_VALUE, &factor, sizeof(factor),
14     0);
15
16 starpu_task_wait_for_all();
17 starpu_data_unregister(vector_handle);
18
19 /* ... display vector ... */
```

Heterogeneity: Device Kernels

Extending a codelet to handle heterogeneous platforms

Heterogeneity: Device Kernels

Extending a codelet to handle heterogeneous platforms

- Multiple kernel implementations for a CPU
 - SSE, AVX, ... optimized kernels

```
1 struct starpu_codelet scal_cl = {  
2     .cpu_func = { scal_cpu_func ,  
3                 scal_sse_cpu_func , scal_avx_cpu_func , NULL } ,  
4     .nbuffers = 1 ,  
5     .modes = { STARPU_RW } ,  
6 };
```

Heterogeneity: Device Kernels

Extending a codelet to handle heterogeneous platforms

- Multiple kernel implementations for a CPU
 - SSE, AVX, ... optimized kernels
- Kernels implementations for accelerator devices
 - OpenCL, NVidia Cuda kernels

```
1 struct starpu_codelet scal_cl = {  
2     .cpu_func = { scal_cpu_func ,  
3         scal_sse_cpu_func , scal_avx_cpu_func , NULL },  
4     .opencl_func = { scal_cpu_opencl , NULL },  
5     .cuda_func = { scal_cpu_cuda , NULL },  
6     .nbuffers = 1,  
7     .modes = { STARPU_RW },  
8 };
```

Writing a Kernel Function for **CUDA**

Writing a Kernel Function for **CUDA**

```
1
2
3
4
5
6
7
8 extern "C" void scal_cuda_func(void *buffers[], void *cl_arg){
9     struct starpu_vector_interface *vector_handle = buffers[0];
10    unsigned n      = STARPU_VECTOR_GET_NX(vector_handle);
11    float *vector = STARPU_VECTOR_GET_PTR(vector_handle);
12    float *ptr_factor = cl_arg;
13
14    ...
15
16
17
18
19 }
```


Writing a Kernel Function for **CUDA**

```
1
2
3
4
5
6
7
8 extern "C" void scal_cuda_func(void *buffers[], void *cl_arg){
9     struct starpu_vector_interface *vector_handle = buffers[0];
10    unsigned n      = STARPU_VECTOR_GET_NX(vector_handle);
11    float *vector = STARPU_VECTOR_GET_PTR(vector_handle);
12    float *ptr_factor = cl_arg;
13
14    unsigned threads_per_block = 64;
15    unsigned nblocks = (n+threads_per_block-1)/threads_per_block;
16
17    ...
18
19 }
```

Writing a Kernel Function for **CUDA**

```
1
2
3
4
5
6
7
8 extern "C" void scal_cuda_func(void *buffers[], void *cl_arg){
9     struct starpu_vector_interface *vector_handle = buffers[0];
10    unsigned n      = STARPU_VECTOR_GET_NX(vector_handle);
11    float *vector = STARPU_VECTOR_GET_PTR(vector_handle);
12    float *ptr_factor = cl_arg;
13
14    unsigned threads_per_block = 64;
15    unsigned nblocks = (n+threads_per_block-1)/threads_per_block;
16
17    vector_mult_cuda<<<nblocks, threads_per_block, 0,
18        starpu_cuda_get_local_stream()>>>(n, vector, *ptr_factor);
19 }
```

Writing a Kernel Function for **CUDA**

```
1 static __global__ void vector_mult_cuda(unsigned n,  
2                                         float *vector, float factor){  
3     unsigned i = blockIdx.x*blockDim.x + threadIdx.x;  
4  
5     ...  
6 }  
7  
8 extern "C" void scal_cuda_func(void *buffers[], void *cl_arg){  
9     struct starpu_vector_interface *vector_handle = buffers[0];  
10    unsigned n = STARPU_VECTOR_GET_NX(vector_handle);  
11    float *vector = STARPU_VECTOR_GET_PTR(vector_handle);  
12    float *ptr_factor = cl_arg;  
13  
14    unsigned threads_per_block = 64;  
15    unsigned nblocks = (n+threads_per_block-1)/threads_per_block;  
16  
17    vector_mult_cuda<<<nblocks, threads_per_block, 0,  
18                    starpu_cuda_get_local_stream()>>>(n, vector, *ptr_factor);  
19 }
```

Writing a Kernel Function for **CUDA**

```
1 static __global__ void vector_mult_cuda(unsigned n,  
2                                         float *vector, float factor){  
3     unsigned i = blockIdx.x*blockDim.x + threadIdx.x;  
4     if (i < n)  
5         vector[i] *= factor;  
6 }  
7  
8 extern "C" void scal_cuda_func(void *buffers[], void *cl_arg){  
9     struct starpu_vector_interface *vector_handle = buffers[0];  
10    unsigned n = STARPU_VECTOR_GET_NX(vector_handle);  
11    float *vector = STARPU_VECTOR_GET_PTR(vector_handle);  
12    float *ptr_factor = cl_arg;  
13  
14    unsigned threads_per_block = 64;  
15    unsigned nblocks = (n+threads_per_block-1)/threads_per_block;  
16  
17    vector_mult_cuda<<<nblocks, threads_per_block, 0,  
18                    starpu_cuda_get_local_stream()>>>(n, vector, *ptr_factor);  
19 }
```